

# Informed Search

## Lecture 4, CMSC 170

John Roy Daradal / Instructor

# Previously on CMSC 170

- Search Problems
- Tree Search
- Uninformed Search
  - Depth-First Search
  - Breadth-First Search
  - Uniform Cost Search

# Today's Topics

- Heuristics
- Greedy Search
- A\* Search
- Graph Search

# Review: Search Problem

- States
- Actions and costs
- Successor function
- Start state
- Goal test

# Review: Search Tree

- **Nodes**: represent plans for reaching states
- Plans have **costs** (sum of action costs)

# Review: Search Algorithm

- Systematically builds a search tree
- Chooses **ordering** of the *fringe*
- **DFS** (stack), **BFS** (queue), **UCS** (PQ)
- **Optimal** = finds least-cost plans

# Review: Uninformed Search

- Explore tree in **systematic** way
- **DFS** (leftmost, deepest), **BFS** (per-layer), **UCS** (least-costly)
- **Limitation**: *unaware* where the **goals** are


# Review: Uniform Cost Search

- **Strategy:** Expand node with **lowest path cost**
- **Good:** UCS is **optimal** and **complete**




# Review: Uniform Cost Search

- **Bad:** explores options in **every direction**
- **Bad:** no info about **goal location**




# Informed Search

## Informed Search



# Informed Search


- *Key Idea*: Use a **heuristic** function
- **Compass**: look at a state and *evaluate* if that state gets you **closer** to the *goal*



# Search Heuristics


- A *function* that **estimates** how close a state is to a goal
- *Input*: state, *output*: score
- Heuristic design **varies** for different search problems

# Search Heuristic



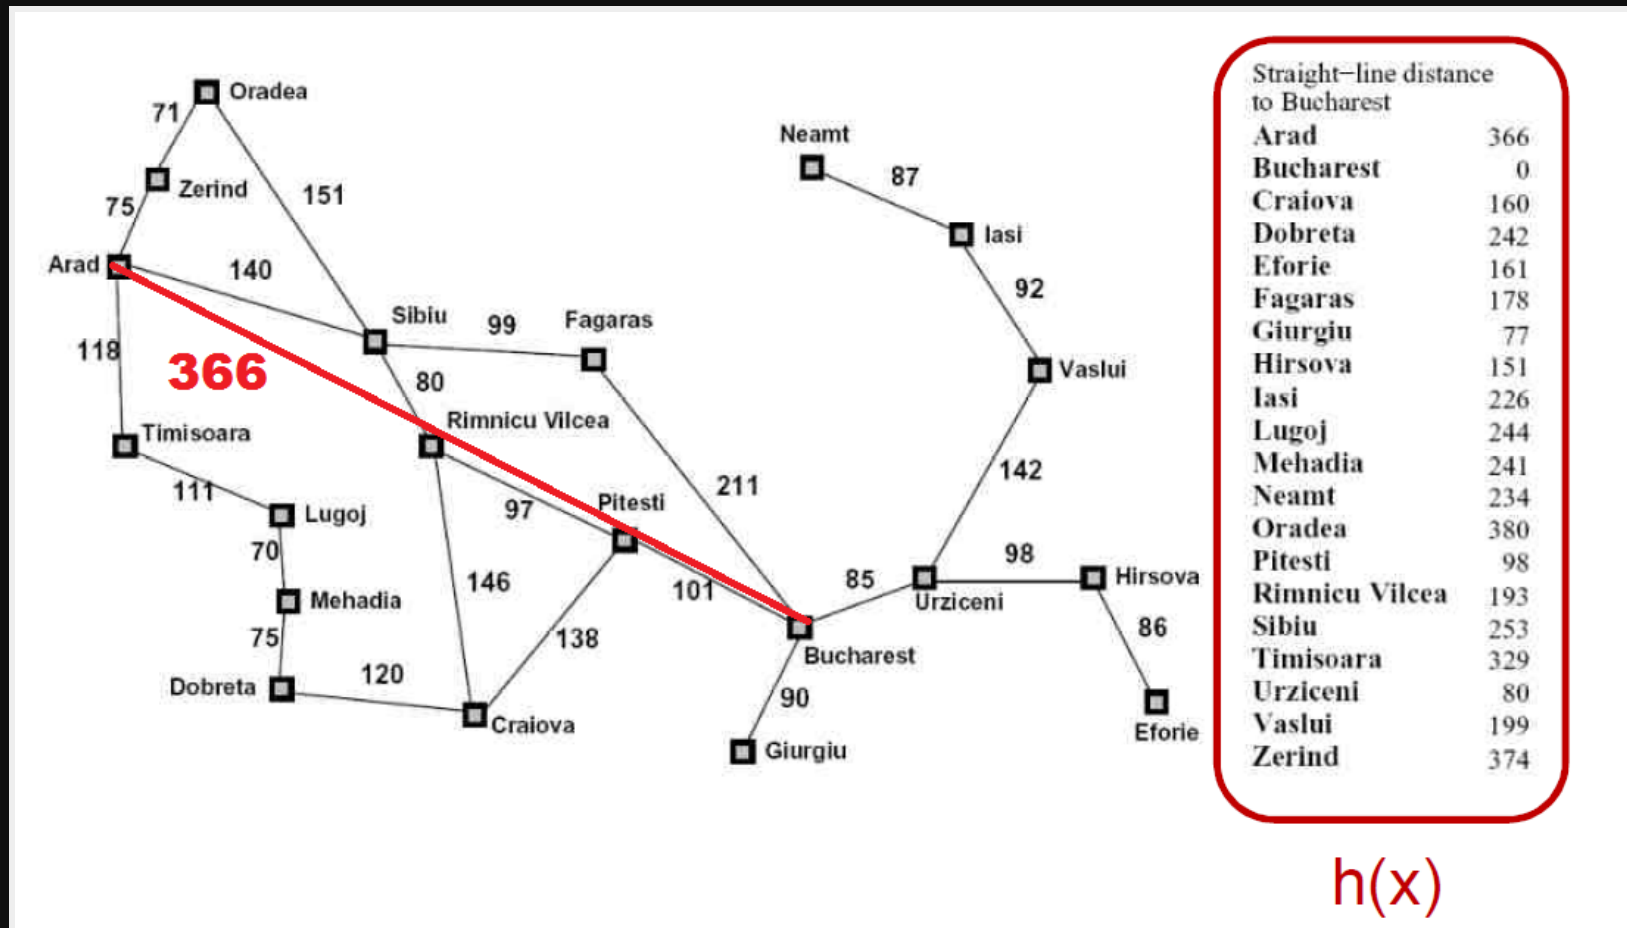
# Example: Pacman

Manhattan distance, Euclidean distance



# Example: Romania Vacation

Straight line distance from city to Bucharest



# Greedy Search





# Greedy Search

- aka *Best-First Search*
- **Strategy**: expand node that **seems closest** to *goal state*, according to **heuristic**

# Greedy Search

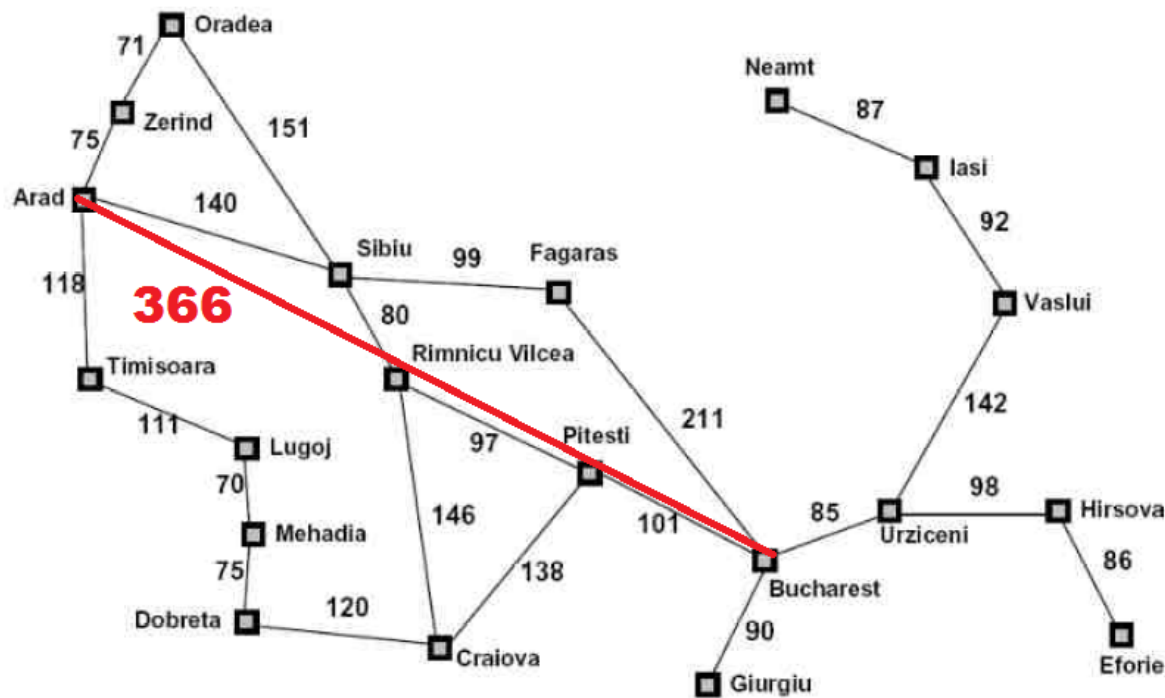
- **Heuristic**: estimates distance to nearest goal for each state
- **Implementation**: fringe = **priority queue**  
(priority: *heuristic value*)

# UCS vs Greedy Search

- *UCS*: **cumulative** cost (sum of action costs)
- *Greedy*: heuristics are **not cumulative**

# Example: Romania Vacation

## UCS vs Greedy Search




Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

$h(x)$

# Greedy Search

What can go wrong?

- **Worst case:** like a badly-guided DFS



# Greedy Search

- Like a **guided DFS**
- *Problem*: when optimal solution needs to do something backwards from heuristic briefly in order to proceed

# Greedy Search

- **Common case**: greedy search takes you *straight* to goal, but with a **suboptimal** plan
- *Search quality* is only as good as **heuristic**
- In general, heuristics are **not perfect**

# Greedy Search Properties

- **Complete?** Yes, if no loops
- **Optimal?** No
- **Time?**  $O(b^m)$ , but good heuristic improves this
- **Space?**  $O(b^m)$ , keeps all nodes in memory



# UCS vs Greedy Search

- *UCS*: **optimal** solution, **no guide**  
(expands in all directions)
- *Greedy*: **guided** search toward goal, mostly **suboptimal** solution

# A\* Search




# Combining UCS and Greedy

- **UCS**: orders by **path cost**, or *backward cost*,  **$g(n)$**
- **Greedy**: orders by **goal proximity**, or *forward cost*,  **$h(n)$**


# Combining UCS and Greedy

- **A\* Search**: orders by the sum  
 $f(n) = g(n) + h(n)$
- *Priority* = **actual path cost + heuristic**

# Example: A\* Search




# Example: A\* Search



- Should A\* stop when we *enqueue* a goal?
- No, stop when we **dequeue** a goal

# Example: A\* Search

Is A\* search optimal?




# What went wrong?


- **Good path** = S-A-G, cost = 4
- **Bad path** = S-G, cost = 5
- **Heuristic** value for A = 6
- actual bad goal cost < estimated good goal cost
- *Estimates* should be **less** than *actual costs*!



# Admissibility



Inadmissible (pessimistic) heuristics break optimality by trapping good plans on the fringe




Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

# Admissible Heuristics

- Let  $h^*(n)$  = true cost to nearest goal
- Heuristic  $h$  is **admissible** (*optimistic*) if  
 **$0 \leq h(n) \leq h^*(n)$**
- estimate  $\leq$  true cost



# Example

Admissible heuristic: Manhattan distance





# UCS vs A\*

Uniform-Cost



**Search Trees**

A\*



**Contours**

# UCS vs A\* Contours

- **UCS**: expands equally in all "directions"
- **A\***: expands mainly toward goal, but trims its bets to ensure optimality

# UCS and A\*

- UCS is A\* search using a *zero heuristic*
- **Zero heuristic** = no clue where the goal is, **worst** kind of heuristic

# Heuristic Quality


*"The worse the heuristic, the lower it is as a lower bound, the more work you're gonna do in order to guarantee optimality."*

# Admissible Heuristics


- *In practice*, designing an **admissible heuristic** is most of the work in using A\* search, to ensure **optimal** results
- *Usually*: solutions to **relaxed problems**, where *new actions* are available



# Creating Heuristics



# Creating Heuristics




Straight-line distance to Bucharest	
Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

$h(x)$

# Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



3	7	1
2	4	5
	8	6

Actions

	1	2
3	4	5
6	7	8


Goal State

**What are possible heuristics  
for 8 Puzzle?**

# Heuristic 1

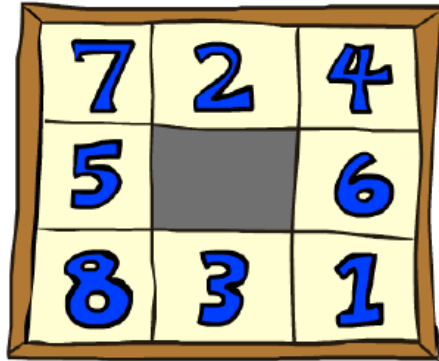
- Number of misplaced tiles
- Loose, extreme relaxation of problem

# Heuristic 1



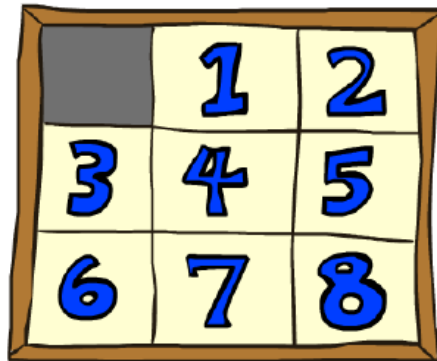
# Heuristic 1

$$h(\text{start}) = 8$$



7	2	4
5		6
8	3	1

Start State



	1	2
3	4	5
6	7	8

Goal State

*Note:* Don't count the free space

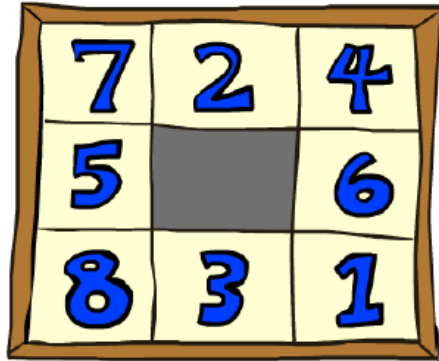
# Heuristic 2

- If we can make any tile slide in any direction (NEWS) at any time, **ignoring other tiles**
- Total Manhattan distance



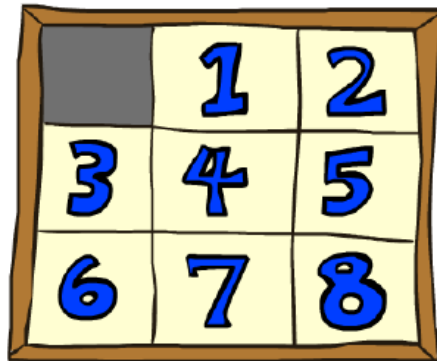
# Heuristic 2

$$h(\text{start}) = 18$$



7	2	4
5		6
8	3	1

Start State



	1	2
3	4	5
6	7	8

Goal State

$$3 + 1 + 2 + 2 + 2 + 3 + 3 + 2$$

# Relaxed Problems

- **8 Puzzle**: a block can move from A to B if  
(1) A is adjacent to B, and (2) B is blank
- Relax condition 2 = heuristic 2
- Relax condition 1,2 = heuristic 1
- Both heuristics are **admissible**

# Heuristic 3

- How about **actual cost** as heuristic?
- Would it be admissible? Yes
- Would we save on nodes expanded? Yes
- What's wrong with it? **Too long** to compute

# A\* Search

- Tradeoff between **quality of estimate** and **work per node**
- As heuristic gets **closer** to *true cost*: **expand fewer** nodes, but do **more work** per node to compute heuristic


# Dominance

- $h_a \geq h_c$  if  $\forall n: h_a(n) \geq h_c(n)$
- *Example:* 8-puzzle, heuristic 2  $\geq$  heuristic 1

# Heuristics

- Max of two or more admissible heuristics is admissible
- $h(n) = \max( h_a(n), h_b(n) )$

# Heuristics




# Recall: Tree Search

```
function TREE-SEARCH(problem, strategy) returns a solution, or failure
  initialize the search tree using the initial state of problem
  loop do
    if there are no candidates for expansion then return failure
    choose a leaf node for expansion according to strategy
    if the node contains a goal state then return the corresponding solution
    else expand the node and add the resulting nodes to the search tree
  end
```




# Recall: Tree Search




Failure to **detect repeated states** can cause **exponentially** more work

# Example

Why shouldn't we bother expanding red nodes?



# Graph Search



# Graph Search

- *Idea*: never expand a state twice
- Tree search + set of expanded states

# Graph Search

- Expand search tree node by node
- Before expanding node, check that it hasn't been expanded before
- If not new, skip it
- If new, expand and add to set

# Graph Search

**function** GRAPH-SEARCH(*problem*, *fringe*) **returns** a solution, or failure

*closed* ← an empty set

*fringe* ← INSERT(MAKE-NODE(INITIAL-STATE[*problem*]), *fringe*)

**loop do**

**if** *fringe* is empty **then return** failure

*node* ← REMOVE-FRONT(*fringe*)

**if** GOAL-TEST(*problem*, STATE[*node*]) **then return** *node*

**if** STATE[*node*] is not in *closed* **then**

        add STATE[*node*] to *closed*

*fringe* ← INSERTALL(EXPAND(*node*, *problem*), *fringe*)


**end**

# Question

- Can graph search destroy completeness? No
- How about optimality? Yes

# Example


State space graph






# Example: Tree Search

Search tree



# Example: Graph Search

Search tree  
using graph search




# Admissible Heuristics

- heuristic cost  $\leq$  actual cost of **path to goal**
- $h(A) \leq$  actual cost from A to G
- **A\* tree search** is *optimal*

# Consistent Heuristics

- heuristic cost  $\leq$  actual cost, **for each arc**
- $h(A) - h(C) \leq$  actual cost from A to C,  
*for each arc A-C*

# Example




# Example

Estimate of cost from:

- A to G =  $h(A) = 4$
- C to G =  $h(C) = 1$
- A to C =  $h(A) - h(C) = 4 - 1 = 3$
- Estimate of cost from A to C **exceeds**  
actual cost = 1

# Consistent Heuristics

Consistent heuristic if  $h(A) = 2$



- $h(A-C) = 2-1 = 1 \leq \text{actual}(A-C) = 1$
- $h(C-G) = 1-0 = 1 \leq \text{actual}(C-G) = 3$

# Consequences of Consistency

- $f$  value along a path never decreases
- $h(A) \leq \text{cost}(A \text{ to } C) + h(C)$
- **A\* graph search** is *optimal*



# A\* Search Properties

- Complete? Yes
- Optimal? Depends on the heuristic

# A\* Search Optimality

Optimal if:

- **A\* tree search**: heuristic is **admissible**
- **A\* graph search**: heuristic is **consistent**

# Heuristics

- Consistency **implies** admissibility
- In general, most **admissible** heuristics tend to be **consistent**, especially if from relaxed problems

# A\* Summary

- Uses both **backward costs** + **estimates of forward cost**
- **Optimal** with admissible / consistent heuristics
- **Heuristic design** is *key*: often use **relaxed** problems

# Informed Search

- Add information about where the goal is
- Find optimal solutions while exploring less of search tree

# Demo: UCS vs Greedy vs A\*

# Search

- Search and Planning
- Completeness vs Optimality
- Uninformed Search (BFS, DFS, ID, UCS)
- Informed Search (Greedy, A\*)
- Admissible vs Consistent Heuristics

# Announcements

- **Assignment 2** available at Courseware
- *Tuesday:* **MP1 - Pacman Search** (Python 2)
- *Next Topic:* Constraint Satisfaction Problems



# References

- *Artificial Intelligence: A Modern Approach, 3rd Edition*, S. Russell and P. Norvig, 2010
- CS 188 Lec 3 slides, Dan Klein, UC Berkeley

# Questions?